

Database Management Systems

Assignment 1

Relational Algebra

Question 1

Consider a database with the following schema:

Person(name, age, gender) // name is a key
Frequents(name, pizzeria) //[name, pizzeria] is a key
Eats(name, pizza) //[name,pizza] is a key
Serves(pizzeria, pizza, price) //[pizzeria,pizza] is a key

Write relational algebra expressions for the following nine queries.

- a. Find all pizzerias frequented by at least one person under the age of 18.
- b. Find the names of all females who eat either mushroom or pepperoni pizza (or both).
- c. Find the names of all females who eat both mushroom and pepperoni pizza.
- d. Find all pizzerias that serve at least one pizza that Amy eats for less than \$10.00.
- e. Find all pizzerias that are frequented by only females or only males.
- f. For each person, find all pizzas the person eats that are not served by any pizzeria the person frequents. Return all such person (name) / pizza pairs.
- g. Find the names of all people who frequent only pizzerias serving at least one pizza they eat.
- h. Find the names of all people who frequent every pizzeria serving at least one pizza they eat.
- i. Find the pizzeria serving the cheapest pepperoni pizza. In the case of ties, return all of the cheapest-pepperoni pizzerias.

Question 2

Consider a schema with two relations, $R(A, B, C)$ and $S(B, C, D)$, where all values are integers. Make no assumptions about keys. Consider the following three relational algebra expressions:

- a. $\pi_{A,D}(R \bowtie \sigma_{B=1}S)$
- b. $\pi_A(\sigma_{B=1}R) \times \pi_D(\sigma_{B=1}S)$
- c. $\pi_{A,D}(\sigma_{B=1}(\pi_{A,B}R) \bowtie \sigma_{B=1}(\pi_{B,D}S))$

Two of the three expressions are equivalent (i.e., produce the same answer on all databases), while one of them can produce a different answer. Which query can produce a different answer? Give the simplest database instance you can think of where a different answer is produced.

Question 3

Consider a relation $R(A, B, C)$ that contains r tuples, and a relation $S(B, C, D)$ that contains s tuples; assume $r > 0$ and $s > 0$. Make no assumptions about keys. For each of the following relational algebra expressions, state in terms of r and s the minimum and maximum number of tuples that could be in the result of the expression.

- a. $R \cup \rho_{S(A,B,C)}S$
- b. $\pi_{A,C}(R \bowtie S)$
- c. $\pi_B R - (\pi_B S - \pi_B R)$
- d. $(R \bowtie S) \bowtie R$
- e. $\sigma_{B>D}S \cup \sigma_{B<C}S$

Question 4

Two more exotic relational algebra operators we didn't cover are the semijoin and antijoin. Semijoin is the same as natural join, except only attributes of the first relation are returned in the result. For example, if we have relations $Enrolled(ID, course)$ and $Student(ID, name)$, and not all students are enrolled in courses, then the query $Enrolled \bowtie Student$ returns the ID and $name$ of all students who are enrolled in at least one course. In the general case, $E1 \bowtie E2$ returns all tuples in the result of expression $E2$ such that there is at least one tuple in the result of $E1$ with matching values for the shared attributes. Antijoin is the converse: $E1 \ltimes E2$ returns all tuples in the result of expression $E2$ such that there are no tuples in the result of $E1$ with matching values for the shared attributes. For example, the query $Enrolled \ltimes Student$ returns the ID and $name$

of all students who are not enrolled in any courses.

Like some other relational operators (e.g., intersection, natural join), semijoin and antijoin are abbreviations - they can be defined in terms of other relational operators. Define $E1 \bowtie E2$ in terms of other relational operators. That is, give an equation $E1 \bowtie E2 = ??$, where ?? on the right-hand side is a relational algebra expression that doesn't use semijoin. Similarly, give an equation $E1 \ltimes E2 = ??$, where ?? on the right-hand side is a relational algebra expression that doesn't use antijoin.

Question 5

Consider a relation $Temp(regionID, name, high, low)$ that records historical high and low temperatures for various regions. Regions have names, but they are identified by $regionID$, which is a key. Consider the following query.

- $T_1(rID, h) = \pi_{regionID, high} Temp$
- $T_2(rID, l) = \pi_{regionID, low} Temp$
- $T_3(regionID) = \pi_{rID}(\sigma_{h < low}(T_1 \bowtie Temp))$
- $T_4(regionID) = \pi_{rID}(\sigma_{l > high}(T_2 \bowtie Temp))$
- $T_5(regionID) = \pi_{regionID} Temp - T_3$
- $T_6(regionID) = \pi_{regionID} Temp - T_4$
- $Result(n) = \pi_{name}(Temp \bowtie (T_5 \cap T_6))$

State in English what is computed as the final Result. The answer can be articulated in a single phrase.