

Chapter 3 Introduction to SQL

Rui Meng
United International College

March 28, 2019

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Value
- Nested Subquery
- View
- Database Modification
- Joined Relations

History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999
 - SQL:2003
 - SQL:2008
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features
 - **Not** all examples here may work on your particular system

Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation
- The domain of values associated with each attribute
- Integrity constraints
- And as we will see later, also other information such as
 - The set of indices to be maintained for each relations
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (e.g., **numeric(3,1)**, allows 44.5 to be stores exactly, but not 444.5 or 0.32).
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- More are covered in Chapter 4.

Creating Table Construct

- An SQL relation is defined using the create table command:

```
create table  $r(A_1D_1, A_2D_2, \dots, A_nD_n,$ 
                (integrity-constraint1),
                ⋮
                (integrity-constraintk))
```

- r is the name of the relation
- each A_i is an attribute name in the schema of relation r
- D_i is the data type of values in the domain of attribute A_i
- Example:

```
create table branch(
    branch_name  varchar(15) not null,
    branch_city  varchar(30),
    assests      integer)
```

Integrity Constraints in Create Table

- **not null**
- **primary key** (A_1, \dots, A_n)
- **foreign key** (A_m, \dots, A_n) **references** r

Example:

```
create table branch(
    branch_name  varchar(15) not null,
    branch_city  varchar(30),
    assests      integer,
    primary key (branch_name),
    foreign key (branch_city) references city )
```

- **primary key** declaration on an attribute automatically ensures not null in SQL-92 onwards, needs to be explicitly stated in SQL-89.

Drop and Alter Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database:

drop table *r*

- The **alter table** command is used to add attributes to an existing relation:

alter table *r* add *A D*

- where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
- All tuples in the relation are assigned **null** as the value for the new attribute.
- The **alter table** command can also be used to drop attributes of a relation:

alter table *r* drop *A*

- where *A* is the name of an attribute of relation *r*.
- Dropping of attributes not supported by many databases (supported by MySQL).

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Value
- Nested Subquery
- View
- Database Modification
- Joined Relations

Basic Query Structure

- A typical SQL query has the form:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

- A_i represents an attribute
 - r_j represents a relation
 - P is a predicate
- This query is equivalent to the following relational algebra expression:

$$\pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

- The result of an SQL query is a relation.

The select Clause

- The **select** clause list the attributes desired in the result of a query
 - corresponds to the projection operation of the relational algebra
- Example: find the names of all branches in the loan relation:

```
select branch_name
from loan
```

- This query is equivalent to the following relational algebra expression:

$$\pi_{branch_name}(loan)$$

- SQL names are case insensitive (i.e., you may use upper- or lower- case letters)
 - E.g., *Branch_Name* \equiv *BRANCH_NAME* \equiv *branch_name*

The select Clause (Cont.)

- SQL **allows duplicates** in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the names of all branches in the loan relations, and remove duplicates.

```
select distinct branch_name
from loan
```

- The keyword **all** specifies that duplicates not be removed.

```
select all branch_name
from loan
```

- Since duplicate retention is the default, we shall not use all in our examples.

The select Clause (Cont.)

- An asterisk in the select clause denotes “all attributes”

```
select *  
from loan
```

- An attribute can be a literal with no from clause, as a temporary constant table

```
select '437'
```

Results is a table with one column and a single row with value “437”. Can give the column a name using:

```
select '437' as FOO
```

- An attribute can be a literal with **from** clause

```
select 'A'  
from loan
```

- Result is a table with one column and N rows (number of tuples in the instructors table), each row with value “A”.

The select Clause (Cont.)

- The **select** clause can contain arithmetic expressions involving the operation, $+$, $-$, $*$, and $/$ and operating on .
- The query:

```
select loan_number, branch_name, amount * 100
from loan
```

would return a relation that is the same as the *loan* relation, except that the value of the attribute *amount* is multiplied by 100.

The where Clause

- The where clause specifies conditions that the result must satisfy.
 - Corresponds to the selection predicate of the relational algebra.
- To find all loan number for loans made at the Perryridge branch with loan amounts greater than \$1200:

```
select loan_number
from loan
where branch_name = 'Perryridge' and amount ≥ 1200
```

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**.
- Comparisons can be applied to results of arithmetic expressions.
- The operands of the logical connectives can be expressions involving the comparison operators $<$, \leq , $>$, \geq , $=$ and \neq .

The where Clause (Cont.)

- SQL includes a between comparison operator.
- Example: Find the loan number of those loans with loan amounts between \$90,000 and \$100,000 (that is, $\geq 90,000$ and $\leq 100,000$):

```

select loan_number
from loan
where amount between 90000 and 100000
    
```

- Similarly, we can use the **not between** comparison operator.

The from Clause

- The **from** clause lists the relations involved in the query.
 - Corresponds to the Cartesian product operation of the relational algebra.

- Example: Find the Cartesian product *borrow* \times *loan*:

select * from *borrow*, *loan*

- Find the name, loan number and loan amount of all customers having a loan at the Perryridge branch.

select *customer_name*, *borrower.loan_number*, *amount*
from *borrower*, *loan*

where *borrower.loan_number* = *loan.loan_number* **and**
branch_name = 'Perryridge'

- Specify the matching condition in **where** clause.

The rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

- The **as** clause can appear in both the **select** and **from** clauses.
 - Rename the attribute name in the query result;
 - Rename relations.
- Example: Find the name, loan number and loan amount of all customers; rename the column name loan_number as loan_id:

```
select customer_name, borrower.loan_number as loan_id, amount
from borrower, loan
where borrower.loan_number = loan.loan_number
```

Tuple Variables

- Tuple variables are defined in the from clause via the use of the as clause.
- Example: Find the customer names and their loan numbers for all customers having a loan at some branch:

```
select customer_name, T.loan_number, S.amount
from borrower as T, loan as S
where T.loan_number = S.loan_number
```

- The identifier T and S , that is used to rename a relation is referred to as a **table alias** or **tuple variable**.
- Example: Find the names of all branches that have greater assets than some branch located in Brooklyn:

```
select distinct T.branch_name
from branch as T, branch as S
where T.assets > S.assets and S.branch_city = 'Brooklyn'
```

String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator “like” uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring.
 - underscore(_). The _ character matches any character.
- Example: Find the names of all customers whose street includes the substring “Main”:

```
select customer_name
from customer
where customer_street like '%Main%'
```

- To match the name “Main%”, we use backslash (\) as the escape character:

```
like 'Main\%' escape '\'
```

Ordering the Display of Tuples

- The **order by** clause causes the tuples in the result of a query to appear in sorted order.
- Example: List in alphabetic order the names of all customers having a loan in Perryridge branch:

```
select distinct customer_name
from borrower, loan
where borrower.loan_number = loan.loan_number and
       branch_name = 'Perryridge'
order by customer_name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
 - Example: **order by** customer_name **desc**
 - Example: **order by** customer_name **desc**, amount **asc**

Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.
- Each duplicate is considered as a distinct tuple.
- **Multiset** versions of some of the relational algebra operators given multiset relations r_1 and r_2 :
 1. $\sigma_\theta(r_1)$: If there are c_1 copies of tuple t_1 in r_1 , and t_1 satisfies selections σ_θ , then there are c_1 copies of t_1 in $\sigma_\theta(r_1)$.
 2. $\pi_A(r_1)$: For each copy of tuple t_1 in r_1 , there is a copy of tuple $\pi_A(t_1)$ in $\pi_A(r_1)$, where $\pi_A(t_1)$ denotes the projection of the single tuple t_i .
 3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2 , there are $c_1 \times c_2$ copies of the tuple $t_1.t_2$ in $r_1 \times r_2$.

Duplicates (Cont.)

- Example: Suppose multiset relations $r_1(A, B)$ and $r_2(C)$ are as follows:

$$r_1 = \{(1, a), (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then $\pi_B(r_1)$ would be $\{(a), (a)\}$, while $\pi_B(r_1) \times r_2$ would be $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$
- SQL duplicate semantics:

```

select  $A_1, A_2, \dots, A_n$ 
from  $r_1, r_2, \dots, r_m$ 
where  $P$ 
    
```

is equivalent to the *multiset* version of the expression:

$$\pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Value
- Nested Subquery
- View
- Database Modification
- Joined Relations

Set Operations

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations \cup , \cap , $-$.
- Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs m times in r and n times in s , then, it occurs:
 - $m + n$ times in r **union all** s
 - $\min(m, n)$ times in r **intersect all** s
 - $\max(0, m - n)$ times in r **except all** s

Set Operations (Cont.)

- Find all customers who have a loan, an account, or both:

```
(select customer_name from depositor)
union
(select customer_name from borrower)
```
- Find all customers who have both a loan and an account:

```
(select customer_name from depositor)
intersect
(select customer_name from borrower)
```
- Find all customers who have an account but no loan:

```
(select customer_name from depositor)
except
(select customer_name from borrower)
```

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Value
- Nested Subquery
- View
- Database Modification
- Joined Relations

Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value:

Aggregate Function	Description
avg	average value
min	minimum value
max	maximum value
sum	sum of values
count	number of values

Aggregate Functions (Cont.)

branch (branch_name, branch_city, assets)

customer (customer_name, customer_street, customer_city)

account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)

- Find the average account balance at the Perryridge branch:

```
select avg(balance)
from account
where branch_name='Perryridge'
```

- Find the number of tuples in the customer relation:

```
select count(*)
from customer
```

- Find the number of depositors in the bank:

```
select count(distinct customer_name)
from depositor
```

Aggregate Functions - Group By

branch (branch_name, branch_city, assets)

customer (customer_name, customer_street, customer_city)

account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)

- Find the number of depositors for each branch:

```
select branch_name, count(distinct customer_name)
from depositor, account
where depositor.account_number = account.account_number
group by branch_name
```

- Note: Attributes in **select** clause outside of aggregate functions must appear in **group by** list.

Aggregate Functions - Having Clause

branch (branch_name, branch_city, assets)

customer (customer_name, customer_street, customer_city)

account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)

- Find the names of all branches where the average account balance is more than \$1,200:

```
select branch_name, avg( balance)
from account
group by branch_name
having avg(balance) > 1200
```

- Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups.

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- **Null Values**
- Nested Subquery
- View
- Database Modification
- Joined Relations

Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes.
- *null* signifies an unknown value or that a value does not exist.
- The predicate **is null** can be used to check for null values.
 - Example: Find all loan number which appear in the *loan* relation with null values for *amount*:


```
select(loan_number)
from loan
where amount is null
```
- The result of any arithmetic expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns *null*.
- However, aggregate functions simply ignores nulls.
 - More on next slide.

Null Values and Aggregates

- Total all loan amounts:

```
select sum (amount)
from loan
```

- Above statement ignores null amounts.
 - Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.
- What if collection has only null values?
 - **count** returns 0
 - All other aggregates return *null*.

Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*.
 - Example: $5 < null$ or $null \neq null$ or $null = null$.
- Three-valued logic using the truth value *unknown*:
 - OR: (*unknown or true*) = *true*,
 (*unknown or false*) = *unknown*,
 (*unknown or unknown*) = *unknown*
 - AND: (*unknown and true*) = *unknown*,
 (*unknown and false*) = *false*,
 (*unknown and unknown*) = *unknown*
 - NOT: (**not** *unknown*) = *unknown*
 - "*P* is **unknown**" evaluates to *true* if predicate *P* evaluates to *unknown*.
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- View
- Database Modification
- Joined Relations

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.
- The nesting can be done in the following SQL query

```

select  $A_1, A_2, \dots, A_n$ 
from  $r_1, r_2, \dots, r_m$ 
where  $P$ 
    
```

as follows:

- A_i can be replaced by a subquery that generates a single value.
- r_i can be replaced by any valid subquery.
- P can be replaced with an expression of the form:
 $B \langle \text{operation} \rangle (\text{subquery})$, where B is an attribute and $\langle \text{operation} \rangle$ to be defined later.

Subquery in where Clause

- A common use of subqueries in the *where* clause is to perform tests:
 - For set membership.
 - For set comparisons.
 - For set cardinality.

Set Membership

branch (branch_name, branch_city, assets)
customer (customer_name, customer_street, customer_city)
account (account_number, branch_name, balance)
loan (loan_number, branch_name, amount)
depositor (customer_name, account_number)
borrower (customer_name, loan_number)

- Find all customers who have both an account and a loan at the bank:

```

select distinct customer_name
from borrower
where customer_name in (select customer_name
                        from depositor )

```
- Find all customers who have a loan at the bank but do not have an account at the bank:

```

select distinct customer_name
from borrower
where customer_name not in (select customer_name
                        from depositor )

```

Set Membership (Cont.)

branch (branch_name, branch_city, assets)
customer (customer_name, customer_street, customer_city)
account (account_number, branch_name, balance)
loan (loan_number, branch_name, amount)
depositor (customer_name, account_number)
borrower (customer_name, loan_number)

- Find all customers who have both an account and a loan at the Perryridge branch:

```

select distinct customer_name
from borrower, loan
where borrower.loan_number = loan.loan_number and
      branch_name='Perryridge' and
      (branch_name,customer_name) in
      (select branch_name, customer_name
       from depositor, account
       where depositor.account_number=account.account_number)
    
```

- Note: Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.

Set Comparison

branch (branch_name, branch_city, assets)
customer (customer_name, customer_street, customer_city)
account (account_number, branch_name, balance)
loan (loan_number, branch_name, amount)
depositor (customer_name, account_number)
borrower (customer_name, loan_number)

- Find all branches that have greater assets than some branch located in Brooklyn:

```

select distinct T.branch_name
from branch as T, branch as S
where T.assets > S.assets and S.branch_city='Brooklyn'
    
```

- Same query using >**some** clause:

```

select branch_name
from branch
where assets >some
    (select assets
     from branch
     where branch_city = 'Brooklyn')
    
```

Definition of “some” Clause

- $F \langle \text{comp} \rangle \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F \langle \text{comp} \rangle t)$.

Where $\langle \text{comp} \rangle$ can be: $<, \leq, >, =, \neq$.

$$(5 \langle \text{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array} \rangle) = \text{true}$$

$$(5 \langle \text{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array} \rangle) = \text{false}$$

$$(5 \langle \text{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array} \rangle) = \text{true}$$

$$(5 \langle \text{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array} \rangle) = \text{true}$$

- Note:

- $(=\text{some}) \equiv \text{in}$.
- However, $(\neq \text{some})$ is not equivalent to **not in**.

Set Comparison (Cont.)

branch (branch_name, branch_city, assets)

customer (customer_name, customer_street, customer_city)

account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)

- Find the names of all branches that have greater assets than all branches located in Brooklyn:

```
select distinct branch_name
```

```
from branch
```

```
where assets > all
```

```
    (select assets
```

```
    from branch
```

```
    where branch_city='Brooklyn')
```

Definition of “all” Clause

- $F \langle \text{comp} \rangle \text{all } r \Leftrightarrow \forall t \in r \text{ such that } (F \langle \text{comp} \rangle t)$. Where $\langle \text{comp} \rangle$ can be: $<, \leq, >, =, \neq$.

$$(5 \langle \text{all} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array} \rangle) = \text{false}$$

$$(5 \langle \text{all} \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array} \rangle) = \text{true}$$

$$(5 \langle \text{all} \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array} \rangle) = \text{false}$$

$$(5 \langle \neq \text{all} \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array} \rangle) = \text{true}$$

- Note:

- $(\neq \text{all}) \equiv \text{not in}$.
- However, $(= \text{all})$ is not equivalent to **in**.

Test for Empty Relations

- The **exists** construct returns the value true if the argument subquery is nonempty.
 - **exists** $r \Leftrightarrow r \neq \emptyset$.
 - **not exists** $r \Leftrightarrow r = \emptyset$.

Use of “exists” Clause

- Find all customers who have an account at all branches located in Brooklyn:

```
select distinct S.customer_name
from depositor as S
where not exists(
    (select branch_name
     from branch
     where branch_city='Brooklyn')
except
    (select R.branch_name
     from depositor as T, account as R
     where T.account_number = R.account_number and
          S.customer_name=T.customer_name))
```

- Note:

- $X - Y = \emptyset \Leftrightarrow X \subseteq Y$.
- Cannot write this query using **=all** and its variants.

Correlated subquery

- A **correlated subquery** (**synchronized subquery**) is a subquery that uses values from the outer query.
- For **each row** processed by the outer query, the subquery is evaluated **once**. It can be inefficient.

Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- Find all customers who have at most one account at the Perryridge branch:

```

select T.customer_name
from depositor as T
where unique (
    select R.customer_name
from account, depositor as R
where T.customer_name = R.customer_name and
       R.account_number = account.account_number
       and account.branch_name = 'Perryridge')
        
```
- Not that **unique** is not widely implemented, MySQL does not support such mechanism.

Test for Absence of Duplicate Tuples (Cont.)

- Find all customers who have at least two accounts at the Perryridge branch:

```

select T.customer_name
from depositor as T
where not unique (
    select R.customer_name
from account, depositor as R
where T.customer_name = R.customer_name and
       R.account_number = account.account_number
       and account.branch_name = 'Perryridge')

```
- Variable from outer level is known as a **correlation variable**.

Subqueries in from Clause

- SQL allows a subquery expression to be used in the **from** clause.
- Example: Find the average account balance of those branches where the average account balance is greater than \$1200.

```
select branch_name, avg_balance
from (select branch_name, avg (balance)
      from account
      group by branch_name )
as branch_avg ( branch_name, avg_balance )
where avg_balance > 1200
```

- Note that we do not need to use the **having** clause, since we compute the temporary (view) relation *branch_avg* in the **from** clause, and the attributes of *branch_avg* can be used directly in the **where** clause.

With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.
- Example: Find all accounts with the maximum balance.

```

with max_balance (value) as
    select max (balance)
    from account
select account_number
from account, max_balance
where account.balance = max_balance.value
    
```

Complex Queries using With Clause

- Find all branches where the total account deposit is greater than the average of the total account deposits at all branches:

```

with branch_total(branch_name, value) as
    select branch_name, sum(balance)
    from account
    group by branch_name
with branch_total_avg(value) as
    select avg (value)
    from branch_total
select branch_name
from branch_total, branch_total_avg
where branch_total.value  $\geq$  branch_total_avg.value
    
```

Subqueries in the Select Clause

- Scalar subquery is one which is used where a single value is expected.
- List all branches along with the total amount deposits in each branch:

```
select branch_name,
      ( select sum(balance)
        from account as T
        where T.branch_name = account.branch_name)
as value
from account
```

- Runtime error if subquery returns more than one result tuple.

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Views
- Database Modification
- Joined Relations

Views

- In some cases, it is not desirable for all users to see the entire **logical model** (that is, all the actual relations stored in the database.)
- Consider a person who needs to know a customers name, loan number and branch name, but has no need to see the loan amount. This person should see a relation described, in SQL, by

```
(select customer_name, borrower.loan_number, branch_name
from borrower, loan
where borrower.loan_number = loan.loan_number )
```

 item A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.

View Definition

- A view is defined using the create view statement which has the form:
create view *v* **as** <query expression>
 where <query expression> is any legal SQL expression. The view name is represented by *v*.
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- When a view is created, the query expression is stored in the database; the expression is substituted into queries using the view.

Query with View

- A view consisting of branches and their customers:

create view *all_customer* **as**

 (**select** *branch_name, customer_name*

from *depositor, account*

where *depositor.account_number = account.account_number*)

union

 (**select** *branch_name, customer_name*

from *borrower, loan*

where *borrower.loan_number = loan.loan_number*)

- Find all customers of the Perryridge branch:

select *customer_name*

from *all_customer*

where *branch_name = 'Perryridge'*

Views Defined Using Other Views

- One view may be used in the expression defining another view.
- A view relation v_1 is said to depend directly on a view relation v_2 if v_2 is used in the expression defining v_1 .
- A view relation v_1 is said to depend on view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2 .

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Views
- Database Modification
- Joined Relations

Deletion

- Delete all account tuples at the Perryridge branch:
delete from *account*
where *branch_name*='Perryridge'
- Delete all accounts at every branch located in the city 'Needham':
delete from *account*
where *branch_name* **in** (**select** *branch_name*
 from *branch*
 where *branch_city* = 'Needham')

Deletion (Cont.)

- Delete the record of all accounts with balances below the average at the bank:

delete from *account*

where *balance* < (**select avg**(*balance*)
from *account*)

- Problem: as we delete tuples from deposit, the average balance changes.
- Solution adopted in SQL:
 1. First, compute **avg** balance and find all tuples to delete.
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Insertion

- Add a new tuple to account:
insert into *account*
 values ('A-9732', 'Perryridge', 1200)
- Or equivalently:
insert into *account* (*branch_name*, *balance*, *account_number*)
 values ('Perryridge', 1200, 'A-9732')
- Add a new tuple to account with balance set to *null*:
insert into *account*
 values ('A-777', 'Perryridge', *null*)

Insertion (Cont.)

- Provide as a gift for all loan customers of the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account:

```
insert into account
```

```
    select loan_number, branch_name, 200
```

```
    from loan
```

```
    where branch_name = 'Perryridge'
```

```
insert into depositor
```

```
    select customer_name, loan.loan_number
```

```
    from loan, borrower
```

```
    where branch_name = 'Perryridge'
```

```
        and loan.loan_number = borrower.loan_number
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation (otherwise queries like **insert into table1 select * from table1** would cause problems)

insert into table1 select * from table1 would cause problems)

Updates

- Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%:
 - Write two update statements:


```
update account
set balance = balance*1.06
where balance > 10000

update account
set balance = balance*1.05
where balance ≤ 10000
```
 - The order is important.
 - Can be done better using the case statement (next slide).

Case Statement for Conditional Updates

- Same query as before: Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%:

```

update account
set balance = case
                when balance ≤ 10000 then balance * 1.05
                else balance * 1.06
                end

```
- The general form of the case statement is as follows:

```

case
  when pred1 then result1
  when pred2 then result2
  ...
  when predn then resultn
  else result0
end

```

Update of a View

- Create a view of all loan data in the loan relation, hiding the amount attribute:

```
create view branch_loan as
    select loan_number, branch_name
from loan
```

- Add a new tuple to *branch_loan*:

```
insert into branch_loan
    values ('L-37', 'Perryridge')
```

This insertion must be represented by the insertion of the tuple ('L-37', 'Perryridge', null) into the *loan* relation.

Update of a View (Cont.)

- Some updates through views are impossible to be translated into “real” updates on the view: **create view *v* as**
select *loan_number, branch_name, amount*
from *loan*
where *branch_name* = 'Perryridge'
insert into *v* **values** ('L-99', 'Downtown', '23')
- Others cannot be translated uniquely:
insert into *all_customer* **values** ('Perryridge', 'John')
 - Have to choose loan or account and create a new loan or account number!
- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation.

create a new loan/account number!

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Views
- Database Modification
- **Joined Relations**

Joined Relations

- **Join operations** take two relations and return as a result another relation. These additional operations are typically used as subquery expressions in the **from** clause.
- **Join condition**: defines which tuples in the two relations match, and what attributes are present in the result of the join.
 - **natural**: retain tuples with same value on common attributes (do not repeat attributes).
 - **on** <predicate>: allows a general predicate over the relations being joined (repeat attributes).
 - **using** (A_1, A_2, \dots, A_n): a form of natural join only requires values to match on specific attributes.

Joined Relations

- **Join type:** defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.
 - **inner join** or **join**, join operations that do not preserve nonmatched tuples.
 - **left outer join:** preserves nonmatched tuples in the first relation.
 - **right outer join:** preserves nonmatched tuples in the second relation.
 - **full outer join:** preserves nonmatched tuples in both relations.

Joined operations

- Relation *course*:

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*:

course_id	pre_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that:
 - prereq* information is missing for CS-315
 - course* information is missing for CS-437

Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

Natural Left Outer Join

- *course* natural left outer join *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

Natural Right Outer Join

- *course natural right outer join prereq*

course_id	pre_id	title	dept_name	credits
BIO-301	BIO-101	Genetics	Biology	4
CS-190	CS-101	Game Design	Comp. Sci.	4
CS-347	CS-101	<i>null</i>	<i>null</i>	<i>null</i>

Natural Full Outer Join

- **course natural full outer join prereq**
 - Not supported in MySQL, use left outer join union right outer join instead.

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

Joined Conditions

- *course* **inner join** *prereq* **on** *course.course_id = prereq.course_id*

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above and a natural join?

Joined Conditions

- *course* **left outer join** *prereq* **on** *course.course_id = prereq.course_id*

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

Joined Conditions (Cont.)

- *course* **full outer join** *prereq* **using** (*course_id*)

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101